

## 1. Problem with current rendering algorithms

Scene → Triangles → Rasterization / Rays → Pixels

It is simulation based.

render images by simulating light interacting with triangles.

Core idea: An image is a mathematical signal, not a geometric process.

replaces simulation with direct mathematical formulas and systems.

Instead of computing light transport every frame, RFR represents the entire visible scene as a compact analytic function and evaluates it directly.

## 2. Concepts

### a. Images as Functions

$$I(\omega) : S^2 \rightarrow \mathbb{R}^3$$

maps every viewing direction  $\omega$  to a color.

Typical rendering algorithms compute  $I$  by physics simulation

My goal is to approximate it using basis functions using radiance functions

## 3. Mathematical model

### a. Basis expansion

The entire image is represented as

$$L(\omega, t, p) = \sum_{i=1}^N a_i(c, t) \Phi_i(\omega)$$

$N$  is a small constant (e.g. 16–64)

$\Phi_i(\omega)$  are fixed at compile time (something like an image template)

Only coefficients  $a_i$  change at compile time

These basis functions are global lighting patterns

## 4. Rendering algorithm

- a. Frame evaluation
  - For each frame
    - Compute coefficients  $a_i$  from camera and time
    - Evaluate basis  $\Phi_i$  for each pixel direction
    - Accumulate image
  - No geometry, no rays, no depth buffers

## 5. System level flow

- a. Pipe line
  - Scene → Radience compiler → coefficient processing + basis set → Radience processing unit → Display image

## 6. Radience compiler

- a. The compiler replaces asset pipelines.
  - 3D scene → Global sampling → basis projection → Coefficient evaluation → Radience program

Something similar to how a video is pre encoded instead of simulated

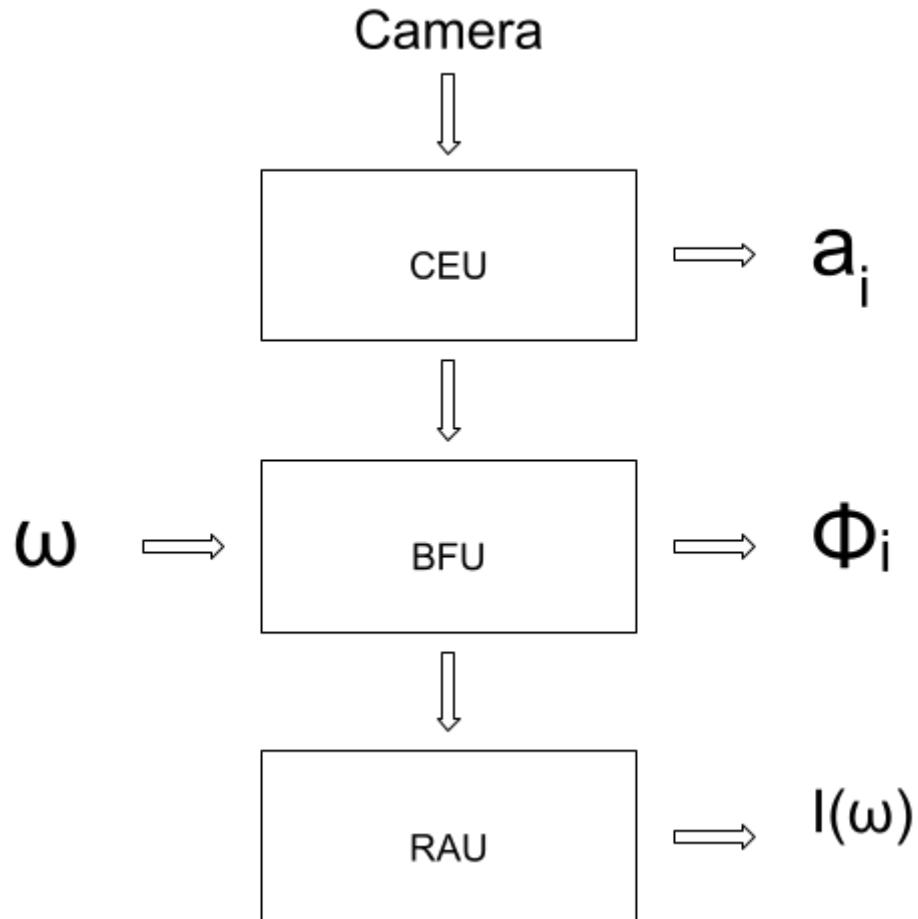
## 7. Runtime pipeline

- a. Per frame execution
  - Camera state → coefficient evaluation unit → basis function unit → radiance accumulation unit → framebuffer

## 8. Radiance processing unit (RPU)

Not a gpu, more like a signal synthesiser  
Each core renders a full frame  
Frame rate can theoretically scale linearly with cores  
1000 cores → 1000 frames per computation cycle

### a. Core structure



## 9. Coefficient Evaluation Unit (CEU)

Computes:

$$a_i = c_0 + c_1x + c_2y + c_3z + c_4\sin(t) + c_5\cos(t)$$

Fully parallel across all basis

## 10. Basis Function Unit (BFU)

Evaluates spherical harmonic or learned analytical bases

Example:

$$\Phi_i(\omega) = Y_l^m(\theta, \phi)$$

## 11. Radiance Accumulation Unit (RAU)

Per pixel:

$$I(\omega) = \sum a_i \Phi_i(\omega)$$

Implemented as a 32 wide vector FMA + reduction tree.

## 12. Instruction set architecture (ISA)

Instruction	Function
CEVAL	Evaluate Coefficients
BEVAL	Evaluate basis
VFMA	Vector multiply accumulate
REDUCE	Sum vector lanes
EMIT	Write pixels

## 13. Pipeline timing

Estimated:

Stage	Cycles
CEU	6
BFU	4
RAU	3
output	1

Total ~14 cycles per frame per core

## 14. Memory Model

No vertex buffers. No textures. No depth buffers

What is there:

storage	size
Basis RAM	KBs
Coefficient RAM	KBs

## 15. Parallel scaling

Each core renders independently  
Performance scales with cores linearly

## 16. Comparison with other GPUs

Feature	GPU	RPU
Geometry	Heavy	None
Memory	TB/s	Minimal
Complexity	High	Lower
Scaling	Sublinear	Linear

## 17. Hybrid system

RPU handles global illumination.

GPU handles:

- UI
- Text
- Particles

## 18. Implementation

- a. Software prototype

Use vectorized CPU code to validate basis accuracy

- b. FPGA prototype

Implement CEU + BFU + RAU pipeline

## 19. Limitations

Issue	Cause
Sharp Edges	Low frequency basis
Dynamic topology	Needs recompilation

## 20. Mathematical derivations

### a. Function space representation

The image is a radiance function:

$$I(\omega) \in L^2(S^2)$$

This is a Hilbert space of square integrable functions on the sphere

$$I(\omega) = \sum \langle I, \Phi_i \rangle \Phi_i(\omega)$$

Where:

$$(f, g) = \int_{S^2} f(\omega) g(\omega) d\omega$$

Therefore, the coefficients are:

$$a_i = \langle I, \Phi_i \rangle$$

Similar to the maths that allows the Fourier Series to represent any signal

### b. Truncated expansion

Using only N terms:

$$I_N(\omega) = \sum_1^N a_i \Phi_i(\omega)$$

The approximation error is:

$$\| I - I_N \|^2 = \sum_{i>N} |a_i|^2$$

The energy outside the first N basis defines the visual error

## 21. Error Bounds and Quality Analysis

### a. Convergence Behavior

Natural images exhibit spectral decay:

$$| a_l | \approx O(l^p) \quad p \approx 2 \text{ to } 4$$

Error decreases rapidly:

$$\| I - I_N \| \approx O(N^{-(p-1)})$$

This ensures high quality with small N

### b. Visual artifact types

Artifact	Cause
Ringing	Gibbs Phenomenon
Blur	Missing high frequency basis
Ghosting	Underfitted coefficients

## 22. Real Basis function definitions

### a. Spherical harmonics

$$\Phi_{l,m}(\theta, \phi) = N_{l,m} P_l^m(\cos\theta) e^{im\phi}$$

$N_{l,m}$  is a normalization constant

$P_l^m(\cos\theta)$  is the associated Legendre polynomial evaluated at  $\cos\theta$

$e^{im\phi}$  is the complex exponential function representing the azimuthal dependence

l is the band

m is the order

### b. Zonal Axis

Axis aligned basis

$$\Phi_l(\theta) = P_l(\cos\theta)$$

Lower cost, good for sky and global lighting

### c. Learned analytic Basis

Train neural network to produce orthogonal functions, then approximate as polynomials for hardware evaluation

## 23. Hybrid GPU integration

### a. Task partitioning

Subsystem	processor
Global lighting	RPU
Geometry edges	GPU
UI/ Text	GPU
Particles	GPU

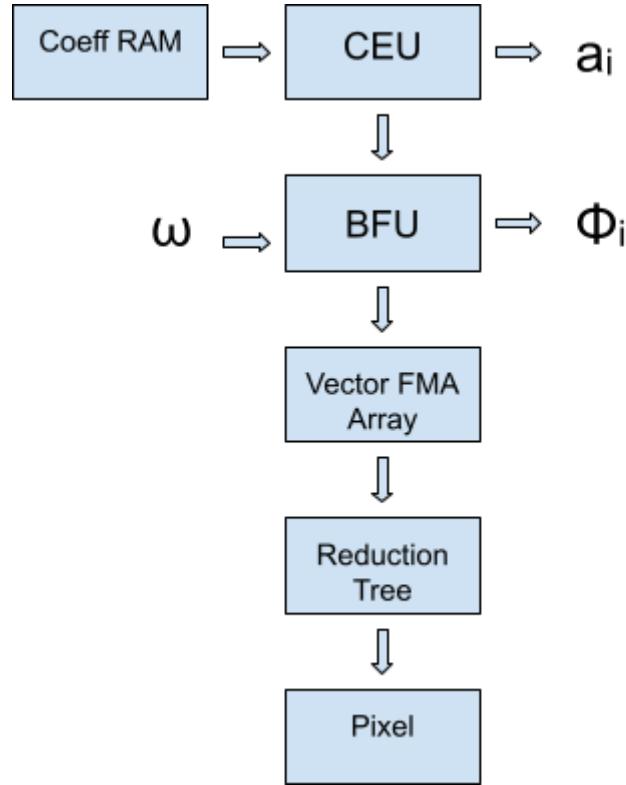
### b. Composite pipeline

RPU → GPU raster pass → final display

RPU acts as a background radiance synthesizer

## 24. Detailed Microarchitecture

### a. Core datapath



### b. CEU pipeline

Regs → FMA → FMA → FMA → TRIG → Output

Fully pipelined, one coefficient per cycle

### c. BFU pipeline

$\omega \rightarrow \sin/\cos \rightarrow \text{polynomial} \rightarrow \text{normalize}$

### d. RAU pipeline

32-wide FMA → 16-wide → 8 → 4 → 2 → 1

Tree reduction for deterministic latency

## 25. Numerical example: toy Scene encoding

### a. Scene definition

Consider a simple scene

- Blue sky
- Green ground
- Bright sun (at  $\theta=30^\circ$ ,  $\phi=0^\circ$ )

True radiance function

$$I(\omega) = 0.5 + 0.3 Y_1^0(\omega) + 1.2 Y_2^1(\omega)$$

### b. Coefficient computation

Using orthonormality:

$$a_i = \int I(\omega) \Phi_i(\omega) d\omega$$

Results:

Basis	$a_i$
$Y_0^0$	0.5
$Y_1^0$	0.3
$Y_2^1$	1.2

All other coefficients approximately 0

### c. Runtime evaluation

Per pixel:

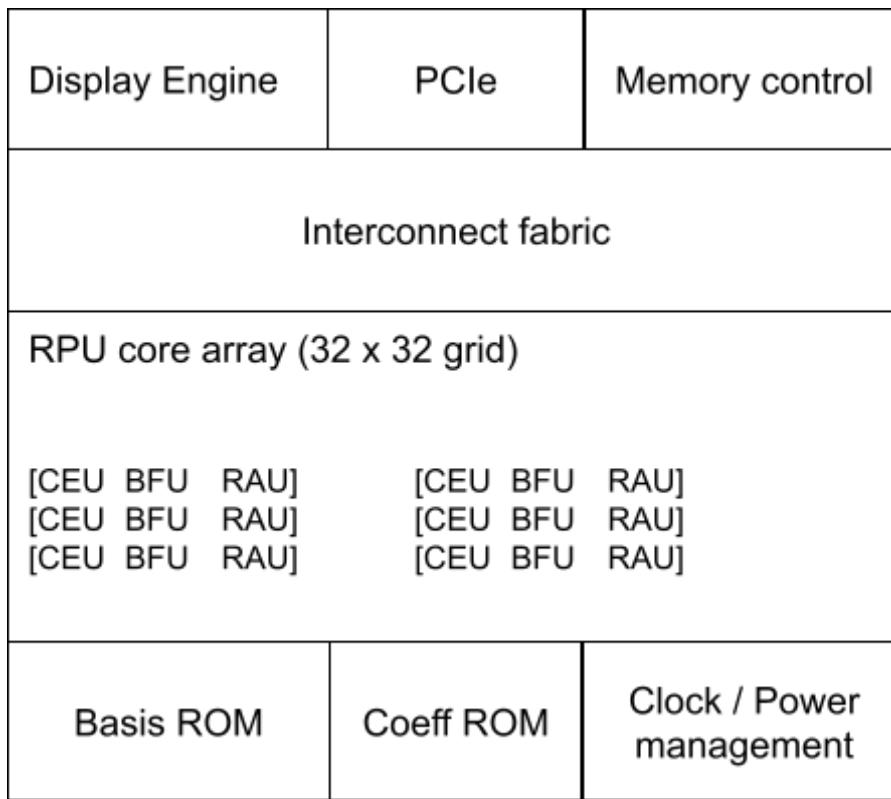
$$I(\omega) \approx 0.5 Y_0^0 + 0.3 Y_1^0(\omega) + 1.2 Y_2^1(\omega)$$

Only 3 FMAs produce the image

## 26. Radiance compiler pseudocode

- a. Global sampling
  - for each probe  $\omega_j$ 
    - $I_j = \text{offline\_path\_trace}(\text{scene}, \omega_j)$
- b. Basis projection
  - for each basis  $\Phi_i$ 
    - $a_i = \sum_j I_j \cdot \Phi_i(\omega_j) \cdot w_j$
- c. Coefficient regression
  - Solve  $\min || a_i(\text{cam}, t) - f_i(\text{cam}, t) ||^2$
  - Store polynomial coefficients
- d. Code estimation
  - Emit CEVAL instructions
  - Emit BEVAL instructions
  - Emit VFMA tree

## 27. Silicon floorplan sketch



Regular layout enables for dense packing

## 28. Performance comparison

### a. Asymptotic Complexity

System	Complexity
Raster	$O(P)$
Ray Tracing	$O(P \log G)$
RFR	$O(1)$

P = pixels

G = geometry

### b. Energy efficiency

Metric	GPU	RPU
Memory BW	TB/s	$\sim 0$
Power	300W	<50W

Based entirely on intuition and estimated guess

## 29. Diagrams

### a. Classical vs RFR pipeline Diagram

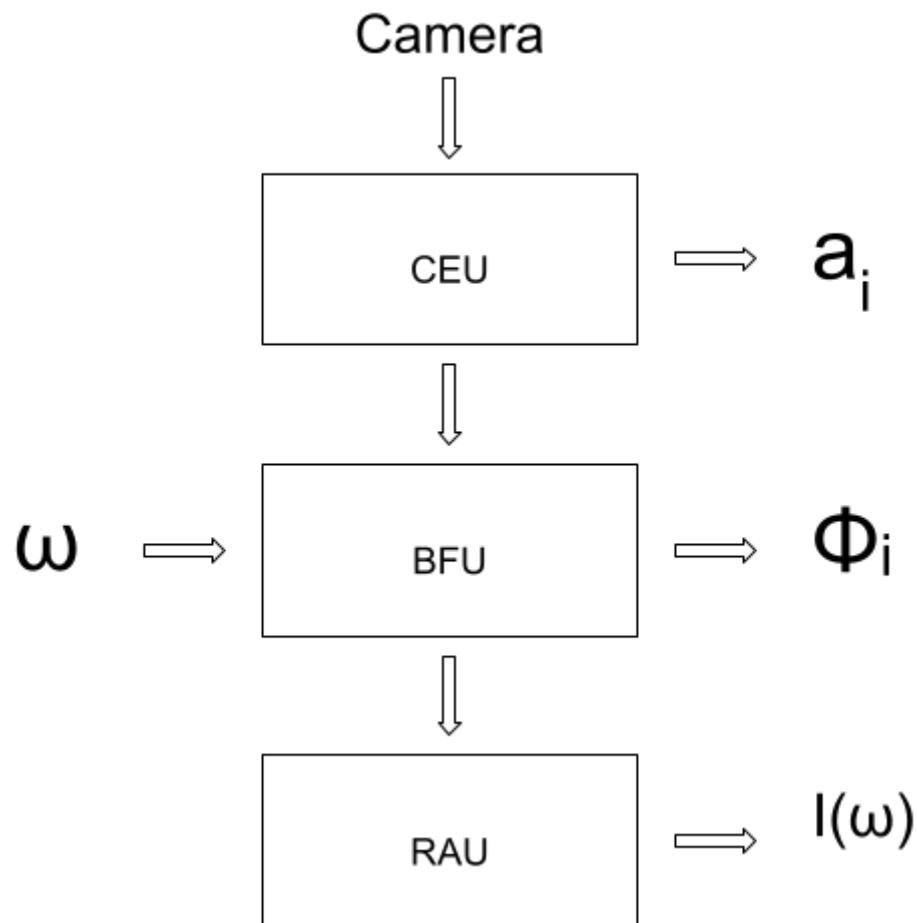
#### i. Classical GPU:

Scene → Triangles → Rasterization → shading → pixels

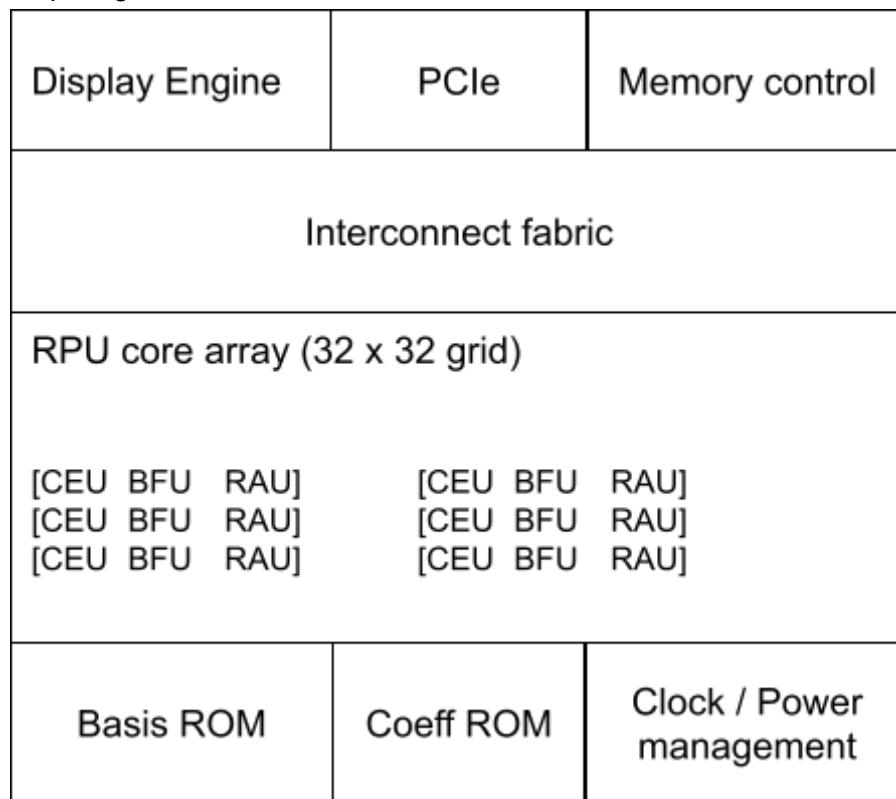
#### ii. RFR:

Scene → Radiance compiler → coefficients + basis → RPU → pixels

#### b. RPU core block



c. Chip diagram



d. Compiler Flow :

3D scene → Global sampling → basis projection → Coefficient regression  
→ Radience Program

### 30. Stress testing (inaccurate values for purpose of example)

- a. High frequency geometry

Problem: sharp edges require many basis terms

Effect

N	Result
16	Blurry edges
32	Acceptable
128	Near perfect

Mitigation: Hybrid GPU edge pass

- b. Rapid topology changes

Explosions or destruction change radiance function suddenly

Issue: coefficients become invalid

Solutions:

- Recompiler coefficients
- Or temporarily fall back to raster GPU

- c. Specular caustics

Highly concentrated light creates high frequency lobes

Requires specialized basis functions (wavelet lobes)

- d. Dynamic lighting

Fast changing lights modify coefficients only

No geometry cost increase

This is a major advantage over raster and RT

- e. Worst case scene

Random noise scene

Energy spectrum is flat → slow convergence

This defines the theoretical limit of RFR

## 31. Conclusion

- Radiance function rendering is not a faster rasterizer
- It is a new computational model for graphics
- By converting rendering from simulation to synthesis, it achieves constant time image generation and perfect parallel scaling
- RFR establishes a new class of processors, Radiance Evaluation Units
- Radiance evaluation converts image synthesis into signal evaluation
- It is theoretically scalable endlessly
- Represents O(1) graphics architecture